

```
In [2]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
```

```
In [3]: import pandas as pd
#imported the dataset given to us
pid=pd.read_csv("C:\\Users\\dhima\\anaconda3\\4rth week\\pima-indians-diabetes.csv")

#made a copy of the original dataset
pid1=pid.copy()
print(pid1)
```

	pregs	plas	pres	skin	test	BMI	pedi	Age	class
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
..
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

[768 rows x 9 columns]

```
In [4]: #pid.columns
#created a list of all the columns present in the dataset
#pid_col=list(pid.columns)
pid2=pid.copy()
#pid_col
#pid_col1=pid_col.copy()
#we do not want to bring changes in the class column so we removed it from the copy
#pid_col1.remove('class')
pid2.drop(["class"], axis = 1, inplace = True)

print(pid2)
```

	pregs	plas	pres	skin	test	BMI	pedi	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
..
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

[768 rows x 8 columns]

Ques 1 Show the performance of K-nearest neighbor (KNN) classifier for different values of K (1, 3, 5, 7, 9,

11, 13, 15, 17, 19, 21)

A. Find confusion matrix (use 'confusion_matrix') for each K.

B. Find the classification accuracy (You can use 'accuracy_score') for each K. Note the value of K for

which the accuracy is high.

```

In [5]: X1 = pid2                                # X denotes the input functions and here class defines
print(X1)
y1 = pid['class']                               # y denotes the output functions
print(y1)
X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1,
                                                         train_size=0.7,
                                                         random_state=42,
                                                         stratify=y1)

print(X1_train)
print(X1_test)
print(y1_train)
print(y1_test)

print(f"Numbers of train instances by class: {np.bincount(y1_train)}")
print(f"Numbers of test instances by class: {np.bincount(y1_test)}")

```

	pregs	plas	pres	skin	test	BMI	pedi	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
..
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

[768 rows x 8 columns]

0	1
1	0
2	1
3	0
4	1
..	..
763	0
764	0
765	0
766	1
767	0

Name: class, Length: 768, dtype: int64

	pregs	plas	pres	skin	test	BMI	pedi	Age
209	7	184	84	33	0	35.5	0.355	41
176	6	85	78	0	0	31.2	0.382	42
147	2	106	64	35	119	30.5	1.400	34
454	2	100	54	28	105	37.8	0.498	24
636	5	104	74	0	0	28.8	0.153	48
..
214	9	112	82	32	175	34.2	0.260	36
113	4	76	62	0	0	34.0	0.391	25
556	1	97	70	40	0	38.1	0.218	30
759	6	190	92	0	0	35.5	0.278	66
107	4	144	58	28	140	29.5	0.287	37

```
[537 rows x 8 columns]
      pregs  plas  pres  skin  test  BMI  pedi  Age
730      3   130   78   23    79  28.4  0.323  34
198      4   109   64   44    99  34.8  0.905  26
24      11   143   94   33   146  36.6  0.254  51
417      4   144   82   32     0  38.5  0.554  37
387      8   105  100   36     0  43.3  0.239  45
..      ...   ...   ...   ...   ...   ...   ...   ...
94       2   142   82   18    64  24.7  0.761  21
437      5   147   75    0     0  29.9  0.434  28
86      13   106   72   54     0  36.6  0.178  45
221      2   158   90    0     0  31.6  0.805  66
19       1   115   70   30    96  34.6  0.529  32
```

```
[231 rows x 8 columns]
209      1
176      0
147      0
454      0
636      0
..
214      1
113      0
556      0
759      1
107      0
```

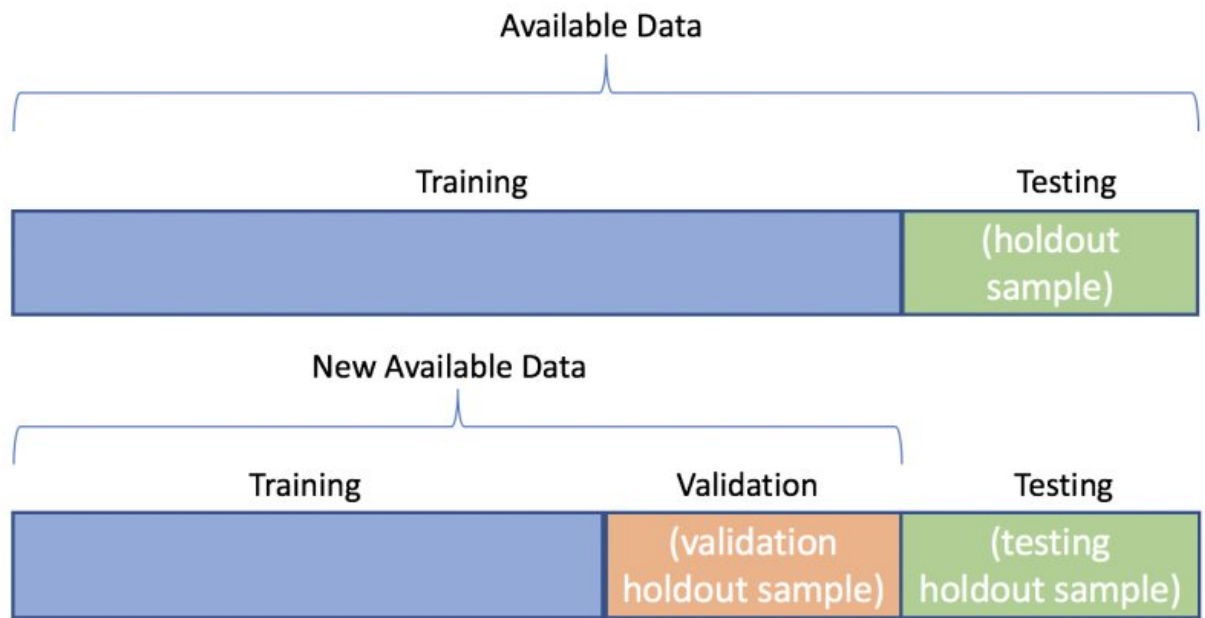
Name: class, Length: 537, dtype: int64

```
730      1
198      1
24       1
417      1
387      1
..
94       0
437      0
86       0
221      1
19       1
```

Name: class, Length: 231, dtype: int64

Numbers of train instances by class: [350 187]

Numbers of test instances by class: [150 81]



```
In [6]: X1_test, X1_val, y1_test, y1_val = train_test_split(X1_test, y1_test,
                                                         train_size=0.5,
                                                         random_state=42,
                                                         stratify=y1_test)

print(X1_train)
print(X1_val)
print(y1_test)
print(y1_val)

print(f"Numbers of test instances by class: {np.bincount(y1_test)}")
print(f"Numbers of validation instances by class: {np.bincount(y1_val)}")
```

	pregs	plas	pres	skin	test	BMI	pedi	Age
209	7	184	84	33	0	35.5	0.355	41
176	6	85	78	0	0	31.2	0.382	42
147	2	106	64	35	119	30.5	1.400	34
454	2	100	54	28	105	37.8	0.498	24
636	5	104	74	0	0	28.8	0.153	48
..
214	9	112	82	32	175	34.2	0.260	36
113	4	76	62	0	0	34.0	0.391	25
556	1	97	70	40	0	38.1	0.218	30
759	6	190	92	0	0	35.5	0.278	66
107	4	144	58	28	140	29.5	0.287	37

[537 rows x 8 columns]

	pregs	plas	pres	skin	test	BMI	pedi	Age
83	0	101	65	28	0	24.6	0.237	22
347	3	116	0	0	0	23.5	0.187	23
52	5	88	66	21	23	24.4	0.342	30
650	1	91	54	25	100	25.2	0.234	23
300	0	167	0	0	0	32.3	0.839	30
..
422	0	102	64	46	78	40.6	0.496	21
580	0	151	90	46	0	42.1	0.371	21
219	5	112	66	0	0	37.8	0.261	41
486	1	139	62	41	480	40.7	0.536	21
112	1	89	76	34	37	31.2	0.192	23

[116 rows x 8 columns]

418	0
235	1
373	0
330	0
64	1
..	..
151	0
535	1
82	0
25	1
637	0

Name: class, Length: 115, dtype: int64

83	0
347	0
52	0
650	0
300	1

```
..
422    0
580    1
219    1
486    0
112    0
Name: class, Length: 116, dtype: int64
Numbers of test instances by class: [75 40]
Numbers of validation instances by class: [75 41]
```

```

In [7]: from sklearn.neighbors import KNeighborsClassifier

neighbors=[1,3,5,7,9,11,13,15,17,19,21]
train_accuracy = np.empty(len(neighbors))
val_accuracy = np.empty(len(neighbors))
acc=[]

# Loop over K values
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X1_train, y1_train)
    print('Predicted Outcomes for neighbours =',k,'are', knn.predict(X1_val))
    print('\n')

    print('Accuracy = ',knn.score(X1_val, y1_val))
    if ((knn.score(X1_val, y1_val))>=0):
        acc.append(knn.score(X1_val, y1_val))
    print('\n')

    matrix = confusion_matrix(y1_val,knn.predict(X1_val))
    print('Confusion Matrix = ',matrix)
    print('\n\n')

# Compute training and test data accuracy
train_accuracy[i] = knn.score(X1_train, y1_train)
val_accuracy[i] = knn.score(X1_val, y1_val)

print(acc)
print('\n')
print('maximum accuracy is =', max(acc)*100)

# Generate plot
plt.plot(neighbors, val_accuracy, label = 'Testing dataset Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training dataset Accuracy')

plt.legend()
plt.xlabel('n_neighbors')
plt.ylabel('Accuracy')
plt.show()

```

```

Predicted Outcomes for neighbours = 1 are [0 1 0 0 1 0 0 1 1 0 0 0 1 0 1 1 1 1
1 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 1 0
0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0
0 0 1 0 0 0 1 1 1 0 1 0 1 0 1 1 1 1 0 1 1 0 0 1 0 0 0 1 1 0 1 0 0 1 0 1 0
0 0 1 1 0]

```

Accuracy = 0.6206896551724138

Confusion Matrix = [[52 23]
[21 20]]

Predicted Outcomes for neighbours = 3 are [0 0 0 0 1 0 0 1 1 0 0 0 1 0 1 1 1 1
 0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0
 0 0 0 0 1 1 0 1 0 0 1 0 0 0 1 1 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 1 0 1 0 0 1 1 0 1 1 1 1 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 0 1 0 1 0
 1 0 1 1 0]

Accuracy = 0.6982758620689655

Confusion Matrix = [[59 16]
 [19 22]]

Predicted Outcomes for neighbours = 5 are [0 1 0 0 1 0 0 0 1 0 0 0 1 0 0 1 1 0
 0 1 0 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1 0
 0 0 0 0 1 1 0 1 0 0 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 1 0
 0 0 1 1 0]

Accuracy = 0.7327586206896551

Confusion Matrix = [[64 11]
 [20 21]]

Predicted Outcomes for neighbours = 7 are [0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 1 1 0
 0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1
 0 0 0 0 1 1 0 1 0 0 1 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0
 0 0 1 1 0]

Accuracy = 0.7413793103448276

Confusion Matrix = [[66 9]
 [21 20]]

Predicted Outcomes for neighbours = 9 are [0 0 0 0 1 0 0 0 1 0 0 0 1 1 0 1 1 0
 0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 1
 0 0 0 0 1 1 0 1 0 0 1 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0
 0 0 1 1 0]

Accuracy = 0.7327586206896551

```
Confusion Matrix = [[64 11]
[20 21]]
```

```
Predicted Outcomes for neighbours = 11 are [0 1 0 0 1 0 0 0 1 0 0 0 1 1 0 1 1 0
0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0
0 0 0 0 1 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0
0 0 0 1 0]
```

```
Accuracy = 0.7068965517241379
```

```
Confusion Matrix = [[64 11]
[23 18]]
```

```
Predicted Outcomes for neighbours = 13 are [0 0 0 0 1 0 0 0 0 0 0 0 1 1 0 1 1 0
0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0
0 0 0 0 1 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0
0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0
0 1 0 1 0]
```

```
Accuracy = 0.7413793103448276
```

```
Confusion Matrix = [[66 9]
[21 20]]
```

```
Predicted Outcomes for neighbours = 15 are [0 0 0 0 1 0 0 0 0 0 0 0 1 1 0 1 1 0
0 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1
0 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0
0 1 0 1 0]
```

```
Accuracy = 0.7672413793103449
```

```
Confusion Matrix = [[68 7]
[20 21]]
```

```
Predicted Outcomes for neighbours = 17 are [0 0 0 0 1 0 0 0 1 0 0 0 1 1 0 1 1 0
0 1 0 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 0 1
0 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0
0 1 0 1 0]
```

Accuracy = 0.7586206896551724

Confusion Matrix = $\begin{bmatrix} 67 & 8 \\ 20 & 21 \end{bmatrix}$

Predicted Outcomes for neighbours = 19 are [0 0 0 0 1 0 0 0 1 0 0 0 1 1 0 1 1 0
0 1 0 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 1
0 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 0
0 1 0 1 0]

Accuracy = 0.7672413793103449

Confusion Matrix = $\begin{bmatrix} 67 & 8 \\ 19 & 22 \end{bmatrix}$

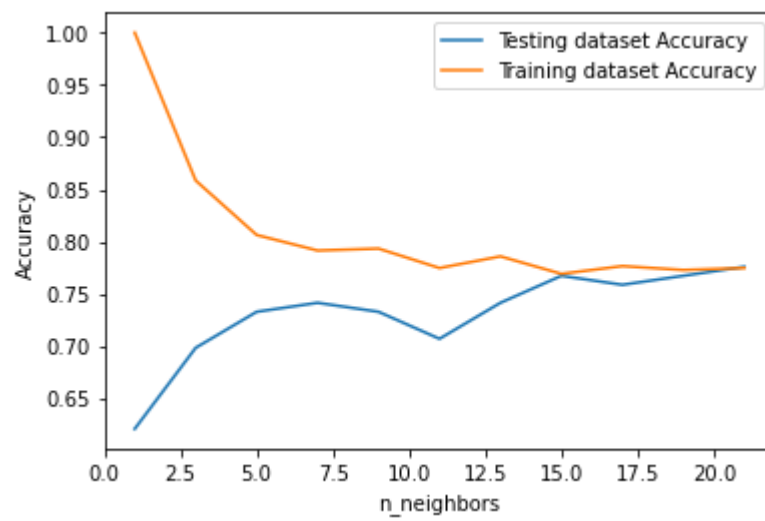
Predicted Outcomes for neighbours = 21 are [0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 1 1 0
0 1 0 1 0 1 1 0 0 1 0 0 0 0 0 0 0 0 0 1
0 0 0 0 1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 1 0
0 1 0 1 0]

Accuracy = 0.7758620689655172

Confusion Matrix = $\begin{bmatrix} 68 & 7 \\ 19 & 22 \end{bmatrix}$

[0.6206896551724138, 0.6982758620689655, 0.7327586206896551, 0.741379310344827
6, 0.7327586206896551, 0.7068965517241379, 0.7413793103448276, 0.76724137931034
49, 0.7586206896551724, 0.7672413793103449, 0.7758620689655172]

maximum accuracy is = 77.58620689655173



```

In [8]: from sklearn.neighbors import KNeighborsClassifier

neighbors=[1,3,5,7,9,11,13,15,17,19,21]
#train_accuracy = np.empty(len(neighbors))
#val_accuracy = np.empty(len(neighbors))
acc=[]

# Loop over K values
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X1_train, y1_train)
    print('Predicted Outcomes for neighbours =',k,'are', knn.predict(X1_test))
    y_pred=knn.predict(X1_test)
    y_pred

    print('Accuracy = ',knn.score(X1_test, y1_test))
    if ((knn.score(X1_test, y1_test))>=0):
        acc.append(knn.score(X1_test, y1_test))
    print('\n')

    matrix = confusion_matrix(y1_test,knn.predict(X1_test))
    print('Confusion Matrix = ',matrix)

    print('\n')

print(acc)
print('\n')
print('maximum accuracy is =', max(acc)*100)

```

```

Predicted Outcomes for neighbours = 1 are [0 1 0 1 1 0 0 0 1 0 0 0 1 0 1 0 1 0
0 0 1 0 1 0 1 0 0 0 1 0 1 0 0 0 1 1 0
0 1 0 0 0 0 1 1 0 1 0 0 0 0 1 0 0 0 0 0 1 0 1 0 0 1 1 1 1 0 0 0 1 0 0 1 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 1 0 0 0 1 1 0 0 0 0 1 0 1 0 0
1 0 0 0]

```

```
Accuracy = 0.7130434782608696
```

```
Confusion Matrix = [[60 15]
[18 22]]
```

```

Predicted Outcomes for neighbours = 3 are [0 1 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0
0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 1 1 0
0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 0 1 1 0 0 0 0 1 0 0 1 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 1 0 0 0 1 1 1 1 0 0 1 0 1 0 0
1 0 0 0]

```

```
Accuracy = 0.7130434782608696
```

```
Confusion Matrix = [[64 11]
[22 18]]
```

```

Predicted Outcomes for neighbours = 5 are [0 1 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0
0 0 1 0 1 0 1 0 0 0 0 1 0 0 0 0 1 1 1
0 0 0 0 0 0 1 0 0 1 0 0 1 0 0 1 0 0 1 0 0 0 0 1 1 1 0 0 0 0 1 1 0 1 0
0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 1 0 0 0 1 1 1 0 0 0 1 0 1 0 0

```

```
1 0 0 0]
Accuracy = 0.7043478260869566
```

```
Confusion Matrix = [[61 14]
[20 20]]
```

```
Predicted Outcomes for neighbours = 7 are [0 1 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0
0 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 1 1 1
0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 1 1 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 1 0 1 0 1 1 1 0 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7391304347826086
```

```
Confusion Matrix = [[63 12]
[18 22]]
```

```
Predicted Outcomes for neighbours = 9 are [0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1
1 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 1 1 1
0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 1 1 0
0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 1 0 1 0 1 1 1 0 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7130434782608696
```

```
Confusion Matrix = [[61 14]
[19 21]]
```

```
Predicted Outcomes for neighbours = 11 are [0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
1 0 1 0 1 0 1 0 0 0 0 1 0 0 0 0 1 0 0
0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 1 1 0 0 0 0 0 1 1 1 1 0
0 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 1 1 1 1 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7043478260869566
```

```
Confusion Matrix = [[62 13]
[21 19]]
```

```
Predicted Outcomes for neighbours = 13 are [0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
1 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 1 0 0
0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 1 1 0
0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 1 1 1 1 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7217391304347827
```

```
Confusion Matrix = [[63 12]
[20 20]]
```

```
Predicted Outcomes for neighbours = 15 are [0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
```

```

1 0 1 0 1 0 1 0 0 0 0 1 1 0 0 0 1 0 0
0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 0 1 0
0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 1 0 1 1 1 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7130434782608696

```

```

Confusion Matrix = [[63 12]
[21 19]]

```

```

Predicted Outcomes for neighbours = 17 are [0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
1 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 1 0 0
0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 0 0 1 0 1 0
0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 1 1 1 1 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7130434782608696

```

```

Confusion Matrix = [[64 11]
[22 18]]

```

```

Predicted Outcomes for neighbours = 19 are [0 1 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0
1 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 1 0 0
0 1 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 1 0 0 1 0 0 0 0 1 1 0 1 0 0 0 0 1 0 1 0
0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 1 1 0 1 0 0 0 1 1 1 1 0 0 1 0 1 0 0
1 0 0 0]
Accuracy = 0.7043478260869566

```

```

Confusion Matrix = [[61 14]
[20 20]]

```

```

Predicted Outcomes for neighbours = 21 are [0 1 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0
1 0 1 0 1 0 1 0 0 0 1 1 1 0 0 0 1 0 0
0 1 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 1 0 0 1 0 0 0 0 1 1 0 1 0 0 0 0 1 0 1 0
0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 1 1 1 1 0 0 1 0 1 0 0
0 0 0 0]
Accuracy = 0.7130434782608696

```

```

Confusion Matrix = [[62 13]
[20 20]]

```

```

[0.7130434782608696, 0.7130434782608696, 0.7043478260869566, 0.739130434782608
6, 0.7130434782608696, 0.7043478260869566, 0.7217391304347827, 0.71304347826086
96, 0.7130434782608696, 0.7043478260869566, 0.7130434782608696]

```

```

maximum accuracy is = 73.91304347826086

```

Ques.2 Build a Bayes classifier with Multi-modal Gaussian distribution (GMM) with Q components (modes) as class

conditional density for each class. Show the performance for different values of Q (2, 4, 8, 16).

Estimate the parameters of the Gaussian Mixture Model (mixture coefficients, mean vectors and covariance matrices) using maximum likelihood method.

A. Find confusion matrix (use 'confusion_matrix') for each Q .

B. Find the classification accuracy (You can use 'accuracy_score') for each Q .

C. Observe the values in the covariance matrix in each case and comment.

D. Compare the results with that obtained using Bayes classifier with unimodal Gaussian distribution

($Q = 1$).

Clustering methods such as K-means have hard boundaries, meaning a data point either belongs to that cluster or it doesn't. On the other hand, clustering methods such as Gaussian Mixture Models (GMM) have soft boundaries, where data points can belong to multiple cluster at the same time but with different degrees of belief. e.g. a data point can have a 60% of belonging to cluster 1, 40% of belonging to cluster 2.

Apart from using it in the context of clustering, one other thing that GMM can be useful for is outlier detection: Due to the fact that we can compute the likelihood of each point being in each cluster, the points with a "relatively" low likelihood (where "relatively" is a threshold that we just determine ourselves) can be labeled as outliers.

This is the exact situation we're in when doing GMM. We have a bunch of data points, we suspect that they came from K different gaussians, but we have no clue which data points came from which gaussian. To solve this problem, we use the EM algorithm. The way it works is that it will start by placing gaussians randomly (generate random mean and variance for the gaussian). Then it will iterate over these two steps until it converges.

E step : With the current means and variances, it's going to figure out the probability of each data point x_i coming from each gaussian. M step : Once it computed these probability assignments it will use these numbers to re-estimate the gaussians' mean and variance to better fit the data points.

That could be a problem for datasets with large number of dimensions (e.g. text data), because with the number of parameters growing roughly as the square of the dimension, it may quickly

become impossible to find a sufficient amount of data to make good inferences. One common way to avoid this problem is to fix the covariance matrix of each component to be diagonal (off-diagonal value will be 0 and will not be updated). To achieve this, we can change the `covariance_type` parameter in scikit-learn's GMM to `diag`.

```

In [9]: X2 = pid2                                # X denotes the input functions and here class defines
print(X2)
y2 = pid['class']                               #y denotes the output functions
print(y2)
X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2,
                                                         train_size=0.7,
                                                         random_state=42,
                                                         stratify=y2)

print(X2_train)
print(X2_test)
print(y2_train)
print(y2_test)

```

	pregs	plas	pres	skin	test	BMI	pedi	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
..
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

[768 rows x 8 columns]

0	1
1	0
2	1
3	0
4	1
..	..
763	0
764	0
765	0
766	1
767	0

Name: class, Length: 768, dtype: int64

	pregs	plas	pres	skin	test	BMI	pedi	Age
209	7	184	84	33	0	35.5	0.355	41
176	6	85	78	0	0	31.2	0.382	42
147	2	106	64	35	119	30.5	1.400	34
454	2	100	54	28	105	37.8	0.498	24
636	5	104	74	0	0	28.8	0.153	48
..
214	9	112	82	32	175	34.2	0.260	36
113	4	76	62	0	0	34.0	0.391	25
556	1	97	70	40	0	38.1	0.218	30
759	6	190	92	0	0	35.5	0.278	66
107	4	144	58	28	140	29.5	0.287	37

[537 rows x 8 columns]

	pregs	plas	pres	skin	test	BMI	pedi	Age
730	3	130	78	23	79	28.4	0.323	34

198	4	109	64	44	99	34.8	0.905	26
24	11	143	94	33	146	36.6	0.254	51
417	4	144	82	32	0	38.5	0.554	37
387	8	105	100	36	0	43.3	0.239	45
..
94	2	142	82	18	64	24.7	0.761	21
437	5	147	75	0	0	29.9	0.434	28
86	13	106	72	54	0	36.6	0.178	45
221	2	158	90	0	0	31.6	0.805	66
19	1	115	70	30	96	34.6	0.529	32

[231 rows x 8 columns]

209 1
176 0
147 0
454 0
636 0

..
214 1
113 0

556 0
759 1
107 0

Name: class, Length: 537, dtype: int64

730 1
198 1
24 1
417 1
387 1

..
94 0
437 0
86 0
221 1
19 1

Name: class, Length: 231, dtype: int64

```
In [10]: import numpy as np
from sklearn.mixture import GaussianMixture
gm = GaussianMixture(n_components=2,covariance_type='full',max_iter=1000,n_init=2)
clf=gm.fit(X2_train, y2_train)
y_pred=gm.predict(X2_test)      #NB classifier assumes that all the features are
print(y_pred)
```


```
[0 0 0 1 1 0 1 1 0 1 0 1 0 1 1 0 0 1 0 1 1 0 0 1 1 0 1 1 1 0 0 1 0 0 0
 0 0 1 0 1 1 0 1 0 0 1 1 1 0 0 0 1 1 0 0 1 1 1 1 0 1 0 1 0 1 0 0 1 1 0 1 1
 1 1 0 0 1 1 0 0 1 0 0 1 1 0 0 0 0 1 1 0 1 0 0 1 1 1 1 0 1 1 1 1 0 0 0 0 1
 0 0 0 1 1 0 1 0 1 1 1 1 1 0 0 1 0 0 1 0 0 0 0 1 0 0 1 1 1 1 0 1 0 0 0 0 1
 1 1 1 1 0 1 1 0 0 0 0 1 0 1 0 1 1 0 1 1 0 1 0 1 1 1 1 1 0 1 1 0 1 0 0 1 0
 1 0 1 1 0 0 0 0 0 1 1 0 1 0 0 0 1 1 0 0 1 0 0 1 0 0 0 1 0 1 1 1 1 1 1 0 1
 0 1 0 0 0 1 1 1 0]
```

```
In [12]: print(clf.bic(pid2))
print(clf.aic(pid2))
```

```
36949.88963138758
36536.59234513744
```

```
In [13]: n_estimators=[1,2,4,8,16]
clfs=[GaussianMixture(n ).fit(pid2) for n in n_estimators]
bics=[clf.bic(pid2) for clf in clfs]
aics=[clf.aic(pid2) for clf in clfs]
print(clfs)
print(bics)
print(aics)
plt.plot(n_estimators, bics, label="BIC")      #Low value of both aic and bic is pr
plt.plot(n_estimators, aics, label="AIC")
plt.legend();
```

```
[GaussianMixture(), GaussianMixture(n_components=2), GaussianMixture(n_componen
ts=4), GaussianMixture(n_components=8), GaussianMixture(n_components=16)]
[45311.82212435199, 36884.02231564019, 36637.73973732567, 36116.433283252394, 3
3669.879740016106]
[45107.49537609349, 36470.72502939005, 35806.50137509224, 34449.31276905238, 30
330.99492188293]
```



```
In [14]: import numpy as np
from sklearn.mixture import GaussianMixture
gm = GaussianMixture(n_components=16,covariance_type='full',max_iter=1000,n_init=
clf=gm.fit(X2_train, y2_train)
y_pred=gm.predict(X2_test)      #NB classifier assumes that all the features are
print(y_pred)

print('\n\n')
print(gm.means_)
print('\n\n')
print(gm.score(pid2, y=None))
print('\n\n')
probs = gm.predict_proba(pid2)
print(probs)
```

```
[ 1  1  0  9  9  1  8  9  1  8  1  5  1  9  6  9  5  1  1  9  6  8  8  1
13  8  8  0  8  8  5 11  1  9 12 11 13  0 13  9  1  9  5  1  9 12  1  8
 8  9  1  1 13  9  9  1  7 13  8  8  9  1  8  7  8  1  8  1  1  9  9  0
 5  5  9  9  7  1  9  8  1  1  9  1  1  9  5 13 13  1 13  5  9  1  8  1
 0  8  8  8  8  7  8  9  9  9  1 11  1 13  8  0  1 11  8  9  1  8  7  8
 9  9  9  9  0  1  9  1 11  8  1  1  1  1  9  7  1  5  8  5  8  6  9  1
 7  1 13  8  8  8  8  9  1  8  8  7  0  1  1  9  7  8  6  8  5  1  8  8
13  8  7  8  8  9  9  8  7  9  8  1  8  1 13  9 11  9  7  9  9  1  1  1
 1  7  8  9 12  9  1  1  1  5  8 13  0  8 11  1  8  1  6  7  9 13  8  8
 8  8  9  8  6  8  1  8  1 11  1  8  9  8  7]
```

```
[[ 3.16955352e+00  1.20855643e+02  6.85147670e+01  2.82706054e+01
  1.69410331e+02  3.35737721e+01  6.00045369e-01  3.16630470e+01]
[ 2.56096203e+00  1.02477466e+02  6.85442609e+01  2.63139563e+01
  7.97271661e+01  3.10383569e+01  4.09326960e-01  2.66711606e+01]
[ 5.20000000e+00  1.65600000e+02  7.24000000e+01  3.26000000e+01
  4.82600000e+02  3.35000000e+01  8.99000000e-01  3.90000000e+01]
[ 1.17491023e+00  9.93445115e+01  6.93672324e+01  1.45426745e+01
  0.00000000e+00  2.35038960e+01  4.51625088e-01  2.38118098e+01]
[ 7.00000000e+00  1.44000000e+02  8.10000000e+01  3.40000000e+01
  4.45000000e+01  2.87750000e+01  4.92750000e-01  5.35000000e+01]
[ 4.45929422e+00  1.17714813e+02  9.94917866e-01  2.99682373e+00
  1.03637278e+00  2.54156772e+01  4.23465426e-01  3.23777928e+01]
[ 3.80009513e+00  1.62905133e+02  7.59982751e+01  3.00003433e+01
  2.81800955e+02  3.67902014e+01  5.92580977e-01  4.05010866e+01]
[ 1.29924246e+00  1.06500047e+02  6.76181878e+01  2.88263568e+01
  1.22479032e+02  3.42512948e+01  6.85807018e-01  2.80835182e+01]
[ 4.97789633e+00  1.25499796e+02  7.51082292e+01  0.00000000e+00
  0.00000000e+00  3.15410869e+01  4.05552620e-01  3.78306967e+01]
[ 4.26388802e+00  1.13186016e+02  7.35388708e+01  3.05317231e+01
  0.00000000e+00  3.34120599e+01  4.12666647e-01  3.50534934e+01]
[ 4.00000000e+00  1.97000000e+02  7.00000000e+01  3.90000000e+01
  7.44000000e+02  3.67000000e+01  2.32900000e+00  3.10000000e+01]
[ 5.70094030e+00  1.60061105e+02  7.52280570e+01  2.95928260e+01
  1.40629935e+02  3.27365471e+01  5.63938687e-01  3.76315845e+01]
[ 6.29989779e+00  1.52200130e+02  7.40005862e+01  3.12997068e+01
  3.51992586e+02  3.39699132e+01  5.73080671e-01  3.90018065e+01]
[ 4.58429871e+00  1.44162373e+02  7.55223564e+01  3.50779894e+01
  2.13550014e+02  3.62254968e+01  6.17786498e-01  3.84020991e+01]
```

```
[2.00000000e+00 1.72000000e+02 7.24000000e+01 4.18000000e+01  
5.43400000e+02 3.87200000e+01 3.88200000e-01 3.32000000e+01]  
[3.25000000e+00 4.42500000e+01 7.60000000e+01 2.35000000e+01  
5.75000000e+00 3.04750000e+01 8.91750000e-01 2.95000000e+01]]
```

-20.868892014013216

```
[[1.86705755e-37 8.59182008e-14 0.00000000e+00 ... 1.03181708e-39  
0.00000000e+00 0.00000000e+00]  
[3.10636735e-43 1.75167068e-06 0.00000000e+00 ... 8.55580542e-46  
0.00000000e+00 0.00000000e+00]  
[2.53194702e-30 5.51066358e-23 0.00000000e+00 ... 8.65024706e-56  
0.00000000e+00 0.00000000e+00]  
...  
[3.28194469e-04 9.99664053e-01 0.00000000e+00 ... 7.18147249e-13  
0.00000000e+00 0.00000000e+00]  
[1.30155539e-39 1.12438212e-20 0.00000000e+00 ... 2.30566179e-43  
0.00000000e+00 0.00000000e+00]  
[2.46643591e-41 9.46508124e-06 0.00000000e+00 ... 1.20813143e-50  
0.00000000e+00 0.00000000e+00]]
```